Two layers.
One binding problem.

A chain layer that can enforce a proposal commitment, and an experience layer that makes the gap between what a proposal said and what it would execute impossible to miss. They were built in parallel. They are not yet merged.

SYSTEM ATLAS | A: CHAIN + INTELLIGENCE · MAIN | B: JUDGE EXPERIENCE · B/JUDGE-MODE | NO COMMON ANCESTOR | 57 B TESTS

01 / COMMIT

Evidence and disclosure, not just calldata

A single commitment over rules, calldata, evidence root, disclosure hash and version. Binding a vote to bytes is table stakes; binding the evidence and the conflicts is the addition.

02 / COMPREHEND

Make the description gap legible

Source-linked claims, a strict validator, a human resolution loop, a recursive decoder and a deterministic report that puts the proposal quote beside the decoded call.

03 / REFUSE

Refuse uncommitted shapes; verify code identity

No hidden multicall child, no unapproved upgrade, no stale version, no code swapped after approval. The executor is the gate; Sentinel is early warning.

SOURCES 4 DOCUMENTS	**EVIDENCE** MERKLE ROOT
INTENT CLAIMS + RULES	**COMMITMENT** HASHES
VOTE A CONTRACTS	**EXECUTOR** A CONTRACTS
REPORT B DECODER	**CONSOLE** B JUDGE UI

REAL / VERIFIED | REFERENCE MODEL | PENDING INTEGRATION | READ TIME · ~10 MIN

> “Voting transparency isn’t informed consent. We made the vote bind to the exact action, and we make the evidence, the conflicts, and the execution verifiable.”
>
> THE BINDING IS TABLE STAKES. COMPREHENSION, EVIDENCE, CONFLICTS AND CODE IDENTITY ARE THE PRODUCT.

01 / THE TRUTH BEFORE THE DIAGRAMS

Two complete halves.
One unmerged project.

READ THIS FIRST

The public repository contains two unrelated histories. The chain layer has checked-in Hardhat artifacts. The experience layer has a working judge console and a TypeScript reference executor. Neither side imports the other.

main 55C636A

Chain + intelligence

Solidity contracts, Hardhat tests, Sentinel, deployments, `intelligence/` modules, and a chain-side console prototype.

11/11 SCENARIO SCORECARD · CHECKED-IN | 44 INTELLIGENCE TESTS | 10 PARITY VECTORS | 11 CONTRACT ABIS

43 ERROR SELECTORS | 15 TEST FILES | NO NPM RUN HERE

× NO COMMON ANCESTOR ≠

b/judge-mode EA78E84

Judge experience

Next.js decision console, evidence/AI/hash pipeline, red-team arena, scorecard, Judge Mode, `/present`, and downloadable HTML/JSON reports.

57 TESTS / 9 FILES | 25 PASS | 9 TEST_REQUIRED | 25/0/9

11 SCENARIOS | 9 ROUTES | REFERENCE EXECUTOR

WHAT THIS MEANS The B app does not call A's contracts, open A's governor, or use A's deployed addresses. Its nine `TEST_REQUIRED` scorecard rows mean **the app is not wired to the chain layer**, not that A produced nothing. A's checked-in 11/11 red-team scorecard is an artifact from A's tree; it was not rerun in the B workspace. The two scorecards must never be merged wholesale.

LOCAL ARCHIPELAGO VISUAL-SITE EXPERIMENT · OUT OF SCOPE | THIS DOCUMENT IS NOT THE SITE

02 / BRANCH TOPOLOGY

Two timelines.
No merge base.

INTEGRATION PENDING

The repo is public. `main` is Person A's chain + intelligence tree. `b/judge-mode` is Person B's judge experience. They do not share history; the merge is a deliberate integration project, not a rebase.

main 55C636A · PUBLIC

Chain + intelligence

- **FDA932C** Contracts, Hardhat suite, Sentinel, deployments
- **DF3A6FC** Direction brief, prior art, parity traps
- **112CE18** Chain-side console, evidence pipeline, scenario reports
- **55C636A** Intelligence modules: evidence, chain, graph, ai-draft, intent-bridge · 44 tests

11/11 SCENARIO SCORECARD · CHECKED-IN | 10 PARITY VECTORS | 43 ERRORS | 11 ABIS

Checked-in artifacts, not rerun in this B workspace. No node_modules or full artifacts/ directory in the worktree.

b/judge-mode EA78E84 · PUBLIC

Judge experience

- **92515A6** Reference executor correctness; audit; status brief
- **792CD0C** Judge Mode, `/present`, HTML/JSON report export · 57 tests
- **EA78E84** Positioning, prior art and Person A direction brief

57 TESTS / 9 FILES | 25 PASS | 9 TEST_REQUIRED | 11 SCENARIOS

The app computes verdicts with the TypeScript reference executor. It imports nothing from A's contracts, intelligence or deployments.

NO COMMON ANCESTOR

A PR cannot merge these histories as-is without a deliberate composition. The default public branch is `main`; the judge experience is on `b/judge-mode`. A local Archipelago visual-site experiment on a newer local `main` is intentionally out of scope.

Three surfaces

CHAIN-SIDE CONSOLE

`main` / `app/` — a prototype backed by a Hardhat runner and the intelligence modules. Goal: prove the arena and the execution logic end-to-end.

A · MAIN | HARDHAT BACKEND

JUDGE-FACING CONSOLE

`b/judge-mode` / `app/` — the decision console with Judge Mode, `/present`, downloadable reports, and a reference executor. Goal: explain and verify the trail for judges.

B · B/JUDGE-MODE | 57 TESTS

VISUAL SITE EXPERIMENT

A separate Archipelago / ecosystem page on a newer local `main`. It is not part of this atlas and is not the product architecture.

OUT OF SCOPE | NOT DOCUMENTED HERE

The merge collisions

These are concrete, not stylistic. Source: chain + intelligence source doc.

BOUNDARY	A · MAIN	B · B/JUDGE-MODE	CONSEQUENCE
Commitment hash	`COMMIT_TYPEHASH` + version hash; ethers	custom tuple: chain, executor, proposalKey, version, actions, intent, evidence root, disclosure; viem	Commitments are not interchangeable
Error map	43 selectors across four contracts	17 `ERROR_DEFINITIONS`	Only a few names/selectors overlap; report mapping is not portable
Executor ABI	`execute( … , (address,uint,bytes)[], Rule[])`	`execute(uint256,uint256,bytes[])`	Adapter required before the console can call the contract
Addresses + manifest	Fixed Hardhat deployments `{ address, abi }`	Deterministic fixture addresses; `{ Name: "0x…" }`	Labels and commitments differ
Scorecards	11 red-team rows, all PASS	34 rows: 25 PASS / 9 TEST_REQUIRED	Merging either file flips the honesty ledger
`packages/shared/hash.ts`	ethers, typehash scheme, kind codes 0–5	viem, custom scheme, kind codes 1–6	One path, two incompatible schemes
Multicall	`multicall(bytes[])` · `0xac9650d8`	custom `multicall((address,bytes)[])` · `0xcaa5c23f`	Encoded bundles are not interchangeable
Graph / voting constants	disclosed ×0.15, undisclosed ×1.2; harmonic voting	disclosed ×0.35, undisclosed ×1.0; B formulas	The same graph yields different scores
`app/`	Chain-side console + Hardhat backend	Judge console + reference executor	Both claim the same path; compose or choose a shell
Intelligence bridge	`intentHash` = canonical JSON; empty multicall bytes inside `intent-bridge.ts`	Reference executor uses its own scheme	Bridge must be corrected, not merely wired

WHAT A SUCCESSFUL MERGE MUST DECIDE Which hash is canonical · which scorecard is the public truth · how the two `app/` shells compose · whether the reference executor remains a fallback · where the shared fixtures live. None of these is a cosmetic refactor.

The whole pipeline,
with the seams visible.

REAL · REFERENCE · PENDING

Read left to right. The first half turns documents into a commitment. The middle is the governed execution path. The last half turns the run into a report a human can inspect. The dotted lines are the unmerged boundary.

SOURCES 4 DOCUMENTS	**EVIDENCE** MERKLE ROOT
CLAIMS PRE-BAKED	**RULES** VALIDATED
COMMITMENT HASHES	**GOVERNOR** A · NOT WIRED
EXECUTOR A · NOT WIRED	**DECODER** B · REAL
REPORT B · REAL	**CONSOLE** B · REFERENCE

Green = present and checked in · amber = reference model · red = specified but not wired across the boundary · blue = A-side contract surface.

LANE 1 · EVIDENCE → INTENT

The four authored documents are hashed into a Merkle root. The pre-baked extractor proposes claims and rules; the validator rejects any claim without a document quote; the reviewer resolves the two ambiguities. The rule set and assumptions produce `intentHash`, `contentHash` and `actionsHash`.

REAL HASHING · BAKED EXTRACTION · SOURCE-LINKED

LANE 2 · COMMITMENT → CHAIN

A's contracts add the append-only registry, signed disclosure, governor vote and executor. The executor recomputes the commitment and refuses anything that does not match the committed rules. Sentinel simulates and flags before execution.

CONTRACTS PRESENT · 11/11 CHECKED-IN RED-TEAM · NOT IMPORTED BY B

LANE 3 · EXECUTION → REPORT

B's reference executor applies its own rule/commitment model to the same demo. The decoder turns calldata into a recursive action tree, maps custom errors, and produces a deterministic report with the quote, the call, the constraint and the evidence.

REFERENCE MODEL · DECODER + REPORT · HTML / JSON

LANE 4 · CONFLICTS → DECISION

B's simulated wallet graph detects direct transfers, common funders and timing. Noisy-OR turns four links into 0.88 RED. The vote brief records the brief hash; the Decision Trail combines claims, COI, voting and the execution hard gate.

SIMULATED WALLETS | REAL MATH | ADVISORY SCORES

SEAM 1 · COMMITMENT

A and B compute different commitments from different preimages. A binds a typehash and version hash; B binds a proposal key, evidence root and disclosure. This is the first thing the integration must reconcile.

SEAM 2 · EXECUTOR

A's `IntentLockExecutor.execute` takes rules and actions; B's reference `execute` takes a three-argument bundle. A deployed call needs an adapter, not a rename.

SEAM 3 · EVIDENCE

A's evidence store and B's evidence manifest are independent. The roots will not match until one store feeds both.

SEAM 4 · TRUTH

A's scorecard says 11/11 scenarios pass. B's scorecard says 9 properties are TEST_REQUIRED because B is not wired to A. Both are true in their own tree; the combined truth is the integration work.

---

04 / CHAIN LAYER · PERSON A

The refusal engine.

MAIN · 55C636A

The chain layer does not decide whether a proposal is wise. It decides whether a transaction is the transaction the vote committed to. The interesting property is refusal: an unapproved selector, a hidden multicall child, a stale version or an implementation whose code hash changed must not execute.

PROPOSALREGISTRY

Append-only, hash-chained proposal versions.

Author-only writes · `prevVersionHash` · `verifyChain` recomputes every hash.

`contracts/ProposalRegistry.sol`

DISCLOSUREREGISTRY

Signed wallet disclosure, bond, challenge and ruling.

EIP-712 nonce · `ecrecover == proposer` · window bounds · verifier-only ruling.

`contracts/DisclosureRegistry.sol`

INTENTGOVERNOR

Lifecycle, vote, informed-brief hash, queueing.

Latest-version equality · disclosure hash match · unchallenged · quorum · for > against.

`contracts/IntentGovernor.sol`

INTENTLOCKEXECUTOR

The gate: state, timelock, commitment, rules, post-conditions.

Recompute commitment · version supersession · rule count · `ActionDecoder.enforce` · external call · allowance post-check.

`contracts/IntentLockExecutor.sol`

ACTIONDECODER

Exact matching for six action kinds and nested multicall.

Target, selector, counterparty, amount, code hash · recursion depth < 2 · 12 stable errors.

`contracts/libs/ActionDecoder.sol`

COMMITMENTHASHER

Pure reference hasher used for parity.

Byte equality with the TypeScript vector set · `COMMIT_TYPEHASH`.

`contracts/CommitmentHasher.sol`

TEST HELPERS

MockUSDC, TreasuryProxy, MaliciousImpl, ReentrantRecipient, VulnerableProxy, FeeModule.

EIP-1967 slots · delegatecall payload · re-entrancy · 10% fee boundary.

`contracts/test-helpers/`

Governor lifecycle



Draft · never assigned | Defeated | Halted | Cancelled

`submitProposal` writes `DisclosurePending` directly; a draft exists only in the registry. `syncFromExecutor` is a manual pull nothing calls — the governor can remain `Queued` while the executor is `Executed` or `Halted`.

Executor gate

01 · RECOMPUTE `actionsHash`, `intentHash`, commitment from chain id, executor and fields.

02 · STATE Must be `Queued`; otherwise `ProposalNotInState`.

03 · TIMELOCK `block.timestamp ≥ eta`; otherwise `TimelockNotExpired`.

04 · VERSION Registry latest must equal the submitted version hash; otherwise `VersionSuperseded`.

05 · RULES Top-level action count must equal rule count; each action is enforced exactly.

06 · EFFECTS FIRST State is written before interactions; a re-entrant recipient cannot execute twice.

07 · POSTCONDITIONS Allowance ceiling only. Balance-drain and native-value cap are not implemented.

HONEST EDGES

· A commitment mismatch surfaces as `ProposalNotInState(Queued, Empty)`, not `CommitmentMismatch`.
· There is no native-value check; `value` is bound in `actionsHash` but no ceiling exists.
· `_postConditionChecks` is an allowance check; the comment about balance decrease is not code.
· `castVote`, `Defeated`, `syncFromExecutor`, `cancelProposal` and `returnBond` have no test in A's checked-in suite.

CODEHASHMISMATCH USES `KECCAK256("")` AS OBSERVED CODE HASH

ONLY `CALL` · NO `DELEGATECALL` PATH

Action kinds

ERC20_TRANSFER · 0

Exact target, selector, recipient and amount ceiling. Calldata length 68.

ERC20_APPROVE · 1

Exact spender, bounded ceiling, explicit unlimited-approval guard.

PARAM_SET · 2

Selector plus first uint256 within `maxAmount`; semantic matching is out of scope.

UPGRADE · 3

Implementation address and `extcodehash` check. Code swapped after the vote reverts.

OWNERSHIP_TRANSFER · 4

Exact new owner; hidden ownership inside a multicall fails the child rule.

MULTICALL · 5

Depth < 2, child count matches, every child consumes its own committed rule.

ERROR SURFACE A emits 43 selectors across four contracts (12 from `ActionDecoder`). B's report maps 17 definitions. Only `CommitmentMismatch` and `ExceedsSpendLimit` currently share a name and selector; `SelectorNotAllowed` is not `SelectorNotAuthorized`, and `UnauthorizedRecipient(address,address)` is not B's one-argument form. This is an integration task, not a UI detail.

Sentinel · early warning, not the gate

SIMULATE

Snapshot state, execute the queued commitment through the real executor, capture, then revert.

DIFF

Balances, allowances, owner/admin, EIP-1967 implementation and code hash, plus routing.

SCAN

Static bytecode scan for `DELEGATECALL`, `CALLCODE`, `SELFDESTRUCT`.

HALT

Guardian submits the report hash; halt only, never execute or edit.

Sentinel observes a sandbox, not the block that will run. It is useful exactly because the executor remains the only guarantee.

CHECKED-IN SCORECARD

11/11

S0–S10 red-team scenarios, all PASS. An artifact from A's tree; not rerun here.

PARITY VECTORS

10

Typehash, actions hash and commitment vectors in A's fixture set.

DEPLOYMENTS

11

Contract ABIs embedded in `deployments/local.json`; chain 31337.

TEST INVENTORY

15

14 Hardhat files + 1 intelligence file. 20 red-team `it()` cases + parity + Sentinel + helpers.

FULL CHAIN + INTELLIGENCE SOURCE DOC →

NO HARDHAT RUN IN THIS WORKSPACE

---

05 / INTELLIGENCE BRIDGE · PERSON A

Turning documents
into a defensible rule set.

MAIN · INTELLIGENCE/

Person A's intelligence modules are the bridge between source documents and the chain. They are real code, but two of them disagree with the executor they are meant to feed. That is the most important integration finding in the project.

EVIDENCE.TS

Content-addressed store, sorted-pair Merkle root, proofs, and a `verifySource` that re-hashes the exact file and bounds the quoted span. An empty tree is an explicit unknown, not a zero hash.

REAL MERKLE

CHAIN.TS

Four sequential stages with thresholds in one `CONFIG`. Missing input yields `UNKNOWN`, never `PASS`. Execution is a hard gate: reverted, flagged or halted → RED; pending with a Sentinel violation → RED; not submitted → UNKNOWN.

REAL UNKNOWN-NEVER-PASS

GRAPH.TS

Simulated wallets, four link types, noisy-OR, disclosed damping and an evidence-bound challenge payload. Like B's graph, it is a risk signal over fake transfers, not a finding of wrongdoing.

SIMULATED DATA REAL MATH

AI-DRAFT.TS

Produces an `IntentSpec`; ambiguous questions stay in `unresolved[]` rather than being guessed. It has heuristic and LLM-shaped paths, but no live model is required to explain the demo.

ADVISORY UNRESOLVED

INTENT-BRIDGE.TS · DIVERGENCE 1

It computes `intentHash = keccak256(canonical JSON)`. The executor computes `keccak256(abi.encode(rules))`. The bridge refuses to run while `unresolved[]` is non-empty, but once it runs, the hash it produces is not the executor's hash.

REAL BUG BREAKS P1

INTENT-BRIDGE.TS · DIVERGENCE 2

It reports `childCount = children.length` for `MULTICALL` but encodes an empty inner `bytes[]`. A bridged multicall can never match an on-chain rule set.

REAL BUG | BREAKS NESTED ENFORCEMENT

THE BRIDGE IS NOT JUST PLUMBING The chain layer and the console agree on the concept of a commitment but not on its preimage, error vocabulary, action encoding or graph constants. Fixing these is the first real integration milestone — before wiring any button to a contract.

UNIT TESTS

44

`it()` cases in one file, eight describe blocks: Merkle, store, claims, COI, voting, execution, end-to-end and bridge.

CONSOLE PROTOTYPE

app/

A chain-side console with `DecisionTrail`, `ViolationExplainer` and a backend that runs the Hardhat arena. It exists on `main`; it is not the judge experience.

CHAIN ARTIFACTS

11

`deployments/local.json` with 11 contract ABIs; 43 errors; 10 parity vectors; 11 stored Sentinel report shapes.

FULL SOURCE DOC → | COUNTED FROM SOURCE; NOT RERUN HERE

---

06 / EXPERIENCE LAYER · PERSON B

Making the decision trail
something a human can inspect.

B/JUDGE-MODE · EA78E84

The console is not a dashboard of opinions. It is a chain of evidence: each route answers one question, every number has a why, and the report is a reproducible artifact rather than a paragraph.

Nine routes

`/`	Verdict and four-stage Decision Trail
`/intent`	Claims, ambiguities, live rule set and hashes

`/timeline`	Hash-chained versions and evidence proofs
`/graph`	Disclosed vs discovered wallets; challenge payload
`/vote`	Risk brief, informed ballot, receipt
`/arena`	Eleven attacks, three layers, expected vs actual
`/arena/[id]`	Violation report; HTML/JSON takeaway
`/scorecard`	25 PASS · 0 FAIL · 9 TEST_REQUIRED
`/present`	Seven-slide offline deck; final slide downloads

Decision Trail

01 · CLAIMS	02 · COI RISK	03 · VOTING	04 · EXECUTION
0.81	**0.88**	**0.71**	**GATE**
GREEN · 6 source-linked claims · 2 ambiguities resolved	RED · 4 undisclosed links · noisy-OR above 0.65	AMBER · 59% informed · 1 FOR voter overlaps a flagged wallet	RED · rejected by the executor · `SelectorNotAuthorized`

Claims and voting are advisory; COI is a risk score where higher is worse; execution is a binary gate with no score. A good average cannot unblock a rejected call.

JUDGE MODE

A guided spine over the same routes, not a separate demo. It starts at the verdict, walks seven stops, and finishes into the deck.

`01 /` `02 /intent` `03 /graph` `04 /vote` `05 /arena` `06 report` `07 /scorecard` `finish → /present`

PRESENTATION + TAKEAWAY

`/present` bypasses the shell and runs seven slides with arrow/space navigation and fullscreen. The last slide, like every report page, downloads a self-contained HTML report or a JSON bundle.

OFFLINE · KEYBOARD · PRINTABLE

57 TESTS / 9 FILES · 25 PASS · 9 TEST_REQUIRED · 25/0/9 · REFERENCE EXECUTOR

---

07 / THE ARTIFACT

One report.
Quote, call, constraint, evidence.

PRIMARY DEMO REPORT

This is what the experience exists to produce. The proposal used implementation language; the queued call is an uncommitted upgrade. The report makes both visible in the same frame and names the violated constraint.

THE PROPOSAL SAID

"We will align the treasury implementation with the program's delivery roadmap during the grant period."

§5 · IMPLEMENTATION NOTE | HUMAN TEXT

≠

THE QUEUED CALL DOES

`upgradeTo()`

Upgrades TreasuryProxy to implementation 0xf7A1…2337

TARGET · TREASURYPROXY | IMPLEMENTATION · 0XF7A1…2337

SelectorNotAuthorized · 0x1420732b

APPROVED COMMITMENT
0xb78b…16d755

SUBMITTED COMMITMENT
0xd9d2…c458a

COMMITTED VS SUBMITTED
3 actions / 1 call · differ

REVERT DATA INCLUDED | SIMULATED TX HASH | DETERMINISTIC CODE, NOT PROSE | DOWNLOAD HTML / JSON

The HTML report repeats the quote, decoded call, every gate, the violated constraint, the evidence hashes and the reference-model disclosure; the JSON is the machine bundle.

---

08 / ATTACK MAP

Eleven attacks.
Three different claims.

EXECUTOR · SENTINEL · DISCLOSURE

Every row carries its layer. “Rejected by the executor,” “simulation detected” and “disclosure challenge filed” are different claims and must never be visually merged.

#	ATTACK	LAYER	EXPECTED VERDICT	WHAT IT PROVES
00	Benign grant	executor	PASSED	Exact committed bundle executes
01	Recipient substitution	executor	UnauthorizedRecipient	Counterparty is in the commitment
02	Amount inflation	executor	ExceedsSpendLimit	Ceilings are per rule
03	Jargon upgrade	executor	SelectorNotAuthorized	No uncommitted function
04	Nested ownership	executor	NestedCallNotAuthorized	Recursive child enforcement
05	Unlimited approve	executor	ExceedsAllowance	Allowance ceiling
06	Delegatecall payload	sentinel	DelegatecallDetected	Simulation flags before execution
07	Re-entrancy	executor	ReentrancyDetected	Re-entrant recipient blocked
08	Early / replay	executor	TimelockNotElapsed	Time and execution flags
09	Stale version	executor	StaleVersion	Superseded commitment refused
10	Undisclosed COI	disclosure	UndisclosedLink	Graph evidence flags the proposal

11/11 MATCH EXPECTED IN B | B RUNS PRECOMPUTED STEPS; NO LIVE EXECUTION | EXPERIENCE SOURCE DOC →

---

09 / STATUS BOARD

What is proven today.
What is only checked in.

HONEST ACCOUNTING

The most useful page in the atlas. A checked-in artifact is not a live run; a reference model is not a contract; an unmerged integration is not a guarantee.

PERSON B · B/JUDGE-MODE

25 PASS

TESTS 57 / 9 FILES · SCORECARD GENERATED

0 FAIL | 9 TEST_REQUIRED | 34 ENTRIES

The 57-test suite, scorecard, production build and desktop/mobile browser QA were run in this workspace on the B code before the docs-only commit. The reference executor still executes nothing on-chain.

PERSON A · MAIN

11/11 scenarios

CHECKED-IN SCORECARD · 44 INTELLIGENCE TESTS

ALL PASS IN ARTIFACT | 15 TEST FILES | NOT RERUN HERE

A's contracts, tests, deployments, vectors and Sentinel reports are present. Neither worktree used for this atlas has `node_modules/` or a full `artifacts/` build, so A's suite was not rerun.

B reference executor + decoder + report + hashes	**Verified in this workspace** · 57 tests, scorecard, build and browser QA
A Solidity contracts + Hardhat suite	**Checked-in artifacts** · 11/11 scenario scorecard, 10 vectors, 43 errors, 11 ABIs; not rerun here
A intelligence modules	**44 unit tests counted** · source present; two bridge divergences found
A/B end-to-end integration	**Pending** · unrelated histories; no imports; 9 TEST_REQUIRED in B
Sentinel halt / on-chain disclosure / guardian	**Fixture or code present** · not a running end-to-end guarantee

NEVER FLATTEN THESE INTO "25/25" The B scorecard is an in-process reference run, not line coverage. The A scorecard is a red-team artifact, not the B app's on-chain proof. The integration scorecard cannot exist until the two trees have one commitment and one executor.

10 / POSITIONING

Why this is not just another Governor.

THE PRECISE CLAIM

Binding a vote to calldata is already solved by Compound and OpenZeppelin Governors. The honest gap is comprehension, evidence, conflicts and code identity at execution. Say that before a judge says it for you.

SYSTEM	IT ALREADY HAS	IT LEAVES OPEN
Compound / OZ Governor	Proposal id = hash of encoded calls	Description is a separate string; no evidence or disclosure commitment
Tally / Tenderly / OZ Defender	Decoding, simulation, state diffs	Advisory display; the executor does not consult it
Zodiac Roles	Standing calldata permissions at execution	A policy, not a per-proposal commitment
IntentLock	Per-proposal commitment + refusal semantics	Evidence and disclosure are in the commitment; the report makes the gap legible

Say these limits out loud

> **A vote that honestly authorizes something bad still executes.** We bind intent to action; we do not judge whether the intent is wise.

> **Self-authored evidence can be internally consistent.** The Merkle root proves the documents did not change after signing; it does not prove they are true.

> **A compromised governor or executor invalidates the guarantee.** Commitment immutability is a security property, not a nicety.

> **The informed-vote receipt raises the cost of voting blind.** It does not stop a delegate voting against their mandate with their eyes open.

The fuller prior-art review and build direction are in person-a-direction.md.

---

11 / DEMO PATH

Seven stops.
Then the deck.

B BRANCH · OFFLINE

Judge Mode in the header walks the same path; `/present` gives the seven-slide close. The app is a local, offline demo; it is not deployed as a site.

01 / · the verdict
02 /intent · resolve two ambiguities
03 /graph · the undisclosed path
04 /vote · read, acknowledge, cast
05 /arena · run 00 and 03
06 report · download the artifact
07 /scorecard · the nine amber gaps
↳ /present · seven slides

JUDGE MODE · 7 STOPS
ALT + ← / →
ESC EXITS
F FULLSCREEN IN THE DECK
HTML + JSON TAKEAWAY

---

12 / READ NEXT + GLOSSARY

The atlas has sources.
Follow them before trusting a claim.

SOURCE OF TRUTH

01 · Chain + intelligence
contracts · tests · Sentinel · intelligence bridge · collision table

02 · Experience layer
routes · Judge Mode · evidence · decoder · graph · scoring · report

GitHub-readable project map
Mermaid system map · branch topology · status

Positioning + prior art

what is already solved · where the value is · limits

Demo script

rehearsable click path · exact expected values

B runbook

quickstart · formulas · fixtures · limitations

Glossary

Intent	The machine-readable rule set the reviewer confirms.
Commitment	The hash the vote binds to: chain, executor, proposal, version, actions, intent, evidence, disclosure.
intentHash	Hash of compiled action rules; sources and confidence do not enter it.
contentHash	Hash of proposal content: rules, assumptions, title and version.
Evidence root	Merkle root over the content-addressed source documents.
Disclosure hash	Commitment to the proposer's signed wallet disclosure.
Noisy-OR	$R = 1 - \Pi(1 - s_j)$; independent conflict links combine into one risk score.
Reference executor	B's TypeScript model of the rules; it executes nothing on-chain.
Hard gate	Execution is binary; it is never averaged with advisory scores.
TEST_REQUIRED	A property that cannot be honestly proven in the current workspace; amber, never green.
Sentinel	Simulation and diff early warning; not the enforcement boundary.
Chain-side console	A's `app/` prototype on `main`, backed by Hardhat and intelligence.
Judge-facing console	B's Next.js app on `b/judge-mode`, with Judge Mode, `/present` and report export.

DOCUMENT, NOT SITE This atlas is a standalone HTML field manual. It has no external assets, no analytics and no network calls. The product is the Next.js console on `b/judge-mode`; this file is the map.

